

Author: Smridhi Mangla

- Reference: https://mesonet.agron.iastate.edu/request/download.phtml?network=CT_ASOS

▼ Data Preparation

```
#load file
from google.colab import drive
drive.mount('/content/drive')
```

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.moun

```
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split
```

```
#load airport weather dataset
df_airport_weather = pd.read_csv("/content/drive/My Drive/Data Science with Python Spring2020
df_airport_weather.shape
```

↳ (9122, 31)

```
#subset the data
df_subset=df_airport_weather[['station','valid','lon','lat','tmpf','dwpf','sknt','relh','drct
df_subset
```

↳

	station	valid	lon	lat	tmpf	dwpf	sknt	relh	drct	p01i	alt:
0	BDL	1/1/2019 0:00	-72.6825	41.9381	M	M	5	M	110	0.01	30.17
1	BDL	1/1/2019 0:05	-72.6825	41.9381	M	M	4	M	120	0.01	30.16
2	BDL	1/1/2019 0:10	-72.6825	41.9381	M	M	4	M	110	0.02	30.16
3	BDL	1/1/2019 0:15	-72.6825	41.9381	M	M	3	M	110	0.03	30.16
4	BDL	1/1/2019 0:20	-72.6825	41.9381	M	M	3	M	100	0.04	30.15
...
9117	BDL	1/30/2019 23:40	-72.6825	41.9381	M	M	18	M	260	M	29.87
9118	BDL	1/30/2019 23:45	-72.6825	41.9381	M	M	19	M	260	M	29.88
9119	BDL	1/30/2019 23:50	-72.6825	41.9381	M	M	23	M	260	M	29.89
9120	BDL	1/30/2019 23:51	-72.6825	41.9381	16	5	23	61.23	270	T	29.89
9121	BDL	1/30/2019 23:55	-72.6825	41.9381	M	M	19	M	260	M	29.89

9122 rows × 13 columns

▼ Data Preprocessing

```
#Cleaning up the tmpf, sknt and dwpf columns 'M'/'T' counts as NA values
df_subset=df_subset.replace('M',np.nan)
df_subset=df_subset.replace('T',np.nan)
df_subset.isna
```

```
<bound method DataFrame.isna of
0      BDL    1/1/2019 0:00 -72.6825  41.9381  ...  0.01  30.17    NaN    NaN
1      BDL    1/1/2019 0:05 -72.6825  41.9381  ...  0.01  30.16    NaN    NaN
2      BDL    1/1/2019 0:10 -72.6825  41.9381  ...  0.02  30.16    NaN    NaN
3      BDL    1/1/2019 0:15 -72.6825  41.9381  ...  0.03  30.16    NaN    NaN
4      BDL    1/1/2019 0:20 -72.6825  41.9381  ...  0.04  30.15    NaN    NaN
...      ...      ...      ...      ...      ...      ...      ...
9117   BDL    1/30/2019 23:40 -72.6825  41.9381  ...   NaN  29.87    NaN    NaN
9118   BDL    1/30/2019 23:45 -72.6825  41.9381  ...   NaN  29.88    NaN    NaN
9119   BDL    1/30/2019 23:50 -72.6825  41.9381  ...   NaN  29.89    NaN    NaN
9120   BDL    1/30/2019 23:51 -72.6825  41.9381  ...   NaN  29.89  1012.3   31
9121   BDL    1/30/2019 23:55 -72.6825  41.9381  ...   NaN  29.89    NaN    NaN
```

[9122 rows x 13 columns]>

```
#Dropping all NA rows
df_subset = df_subset.dropna()
df_subset
```



	station	valid	lon	lat	tmpf	dwpf	sknt	relh	drct	p01i	alt:
207	BDL	1/1/2019 15:51	-72.6825	41.9381	53.1	41	17	63.38	320	0	29.61
220	BDL	1/1/2019 16:51	-72.6825	41.9381	51.1	33.1	19	49.98	310	0	29.74
232	BDL	1/1/2019 17:51	-72.6825	41.9381	48.9	30.9	20	49.63	310	0	29.71
244	BDL	1/1/2019 18:51	-72.6825	41.9381	46	30	17	53.37	330	0	29.81
255	BDL	1/1/2019 19:51	-72.6825	41.9381	43	28.9	27	57.21	300	0	29.81
...
8141	BDL	1/27/2019 19:51	-72.6825	41.9381	44.1	27	14	50.73	190	0	29.91
8297	BDL	1/28/2019 7:51	-72.6825	41.9381	24.1	10	13	54.36	340	0	30.01
8950	BDL	1/30/2019 10:51	-72.6825	41.9381	25	15.1	13	65.58	330	0	29.81
9053	BDL	1/30/2019 18:51	-72.6825	41.9381	27	7	14	42.11	230	0	29.71
9077	BDL	1/30/2019 20:51	-72.6825	41.9381	28	7	8	40.41	190	0	29.71

125 rows × 13 columns

```
#Converting date using proper format
```

```
df_subset.dtypes
```

```
df_subset['valid']= pd.to_datetime(df_subset['valid'])
```

```
df_subset['valid']
```

```
207    2019-01-01 15:51:00
220    2019-01-01 16:51:00
232    2019-01-01 17:51:00
244    2019-01-01 18:51:00
255    2019-01-01 19:51:00
```

...

```
8141    2019-01-27 19:51:00
8297    2019-01-28 07:51:00
8950    2019-01-30 10:51:00
9053    2019-01-30 18:51:00
9077    2019-01-30 20:51:00
```

```
Name: valid, Length: 125, dtype: datetime64[ns]
```

```
#Creating columns for MONTH, DAY and YEAR
```

```
df_subset['day'] = pd.to_datetime(df_subset['valid']).dt.day
```

```
df_subset['month'] = pd.to_datetime(df_subset['valid']).dt.month
```

```
df_subset['year'] = pd.to_datetime(df_subset['valid']).dt.year
```

```
df_subset
```

	station	valid	lon	lat	tmpf	dwpf	sknt	relh	drct	p01i	alti	ms
207	BDL	2019-01-01 15:51:00	-72.6825	41.9381	53.1	41	17	63.38	320	0	29.68	1006
220	BDL	2019-01-01 16:51:00	-72.6825	41.9381	51.1	33.1	19	49.98	310	0	29.74	1007
232	BDL	2019-01-01 17:51:00	-72.6825	41.9381	48.9	30.9	20	49.63	310	0	29.77	1008
244	BDL	2019-01-01 18:51:00	-72.6825	41.9381	46	30	17	53.37	330	0	29.81	1009
255	BDL	2019-01-01 19:51:00	-72.6825	41.9381	43	28.9	27	57.21	300	0	29.88	1010
...

#Creating a new column called MonthDay in “MM-DD” format

```
df_subset['MonthDay'] = df_subset['month'].map(str) + '-' + df_subset['day'].map(str)
df_subset
```



	station	valid	lon	lat	tmpf	dwpf	sknt	relh	drct	p01i	alti	ms
207	BDL	2019-01-01 15:51:00	-72.6825	41.9381	53.1	41	17	63.38	320	0	29.68	1006
220	BDL	2019-01-01 16:51:00	-72.6825	41.9381	51.1	33.1	19	49.98	310	0	29.74	1007
232	BDL	2019-01-01 17:51:00	-72.6825	41.9381	48.9	30.9	20	49.63	310	0	29.77	1008
244	BDL	2019-01-01 18:51:00	-72.6825	41.9381	46	30	17	53.37	330	0	29.81	1009
255	BDL	2019-01-01 19:51:00	-72.6825	41.9381	43	28.9	27	57.21	300	0	29.88	1010
...

▼ Data Recoding

```
# Creating a binary flag column for rows/observations that are cold (below 32 degF)
#If below 32 degF, call it "COLD";
#otherwise, call it "HOT"
df_subset['tmpf'] = df_subset['tmpf'].astype(float)
df_subset['tmpf_b'] = df_subset['tmpf'].apply(lambda x: 'COLD' if x < 32 else 'HOT')
df_subset['tmpf_b']
```

```
↳ 207      HOT
   220      HOT
   232      HOT
   244      HOT
   255      HOT
   ...
  8141      HOT
  8297     COLD
  8950     COLD
  9053     COLD
  9077     COLD
   Name: tmpf_b, Length: 125, dtype: object
```

```
#Creating a binary flag column for rows/observations are windy.
#If above 20 miles per hour, call it "WINDY",
#otherwise call it "CALM"
df_subset['sknt'] = df_subset['sknt'].astype(float)
df_subset['sknt_mph'] = df_subset['sknt']*1.15078 #converting knots to miles per hour
df_subset['sknt_b'] = df_subset['sknt_mph'].apply(lambda x: 'CALM' if x < 20 else 'WINDY')
df_subset['sknt_b']
```

```
↳ 207      CALM
   220     WINDY
   232     WINDY
   244      CALM
   255     WINDY
   ...
  8141     CALM
  8297     CALM
  8950     CALM
  9053     CALM
  9077     CALM
   Name: sknt_b, Length: 125, dtype: object
```


▼ Aggregation

```
#show the counts of each weather type.
freq_table = pd.crosstab(index=df_subset['sknt_b'], columns=df_subset['tmpf_b'])
freq_table
```

```
↳
```

tmpf_b	COLD	HOT
sknt_b		
CALM	56	44
WINDY	18	7

```
#A table that shows the average wind speed and average temperature per day.  
df_avg=df_subset.groupby(['day','month','year']).agg({'tmpf':'mean','sknt':'mean'})  
df_avg
```

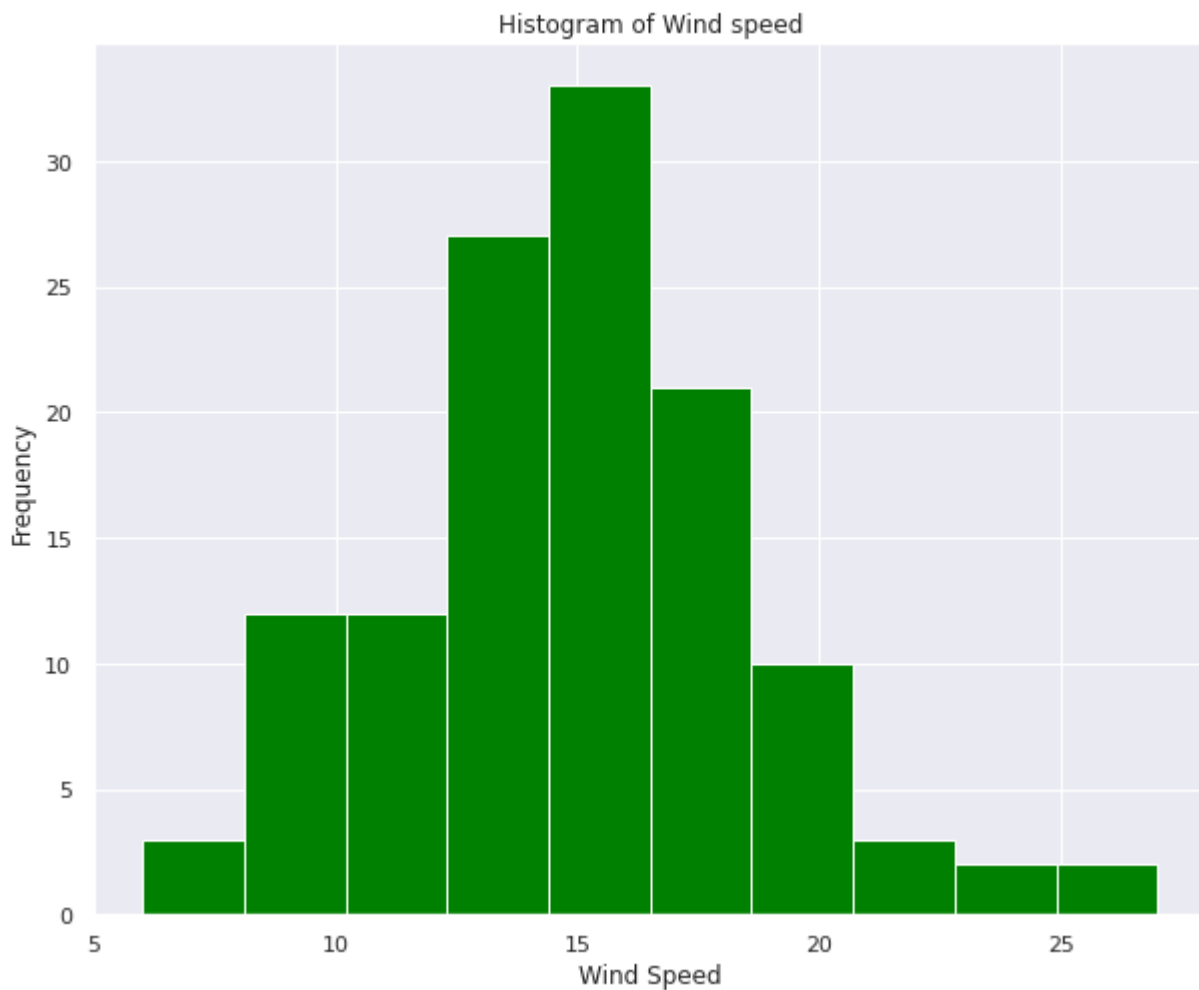


			tmpf	sknt
day	month	year		
1	1	2019	44.862500	17.875000
2	1	2019	37.000000	10.000000
3	1	2019	41.666667	15.666667
4	1	2019	39.000000	10.000000
6	1	2019	39.040000	14.800000
7	1	2019	25.683333	13.833333
9	1	2019	38.828571	15.142857
10	1	2019	30.176471	16.000000
11	1	2019	22.700000	13.900000
12	1	2019	17.380000	11.400000
17	1	2019	29.050000	13.000000
20	1	2019	18.550000	15.500000
21	1	2019	0.526316	16.684211
22	1	2019	6.655556	14.333333
24	1	2019	50.620000	14.800000
25	1	2019	34.033333	15.888889
26	1	2019	29.450000	9.000000
27	1	2019	42.700000	13.666667
28	1	2019	24.100000	13.000000
29	1	2019	26.666667	11.666667

▼ Data Visualization

```
#histogram plot of the wind speed
import matplotlib.pyplot as plt
df_subset.hist('sknt', color='green')
plt.title("Histogram of Wind speed")
plt.xlabel('Wind Speed')
plt.ylabel('Frequency')
```

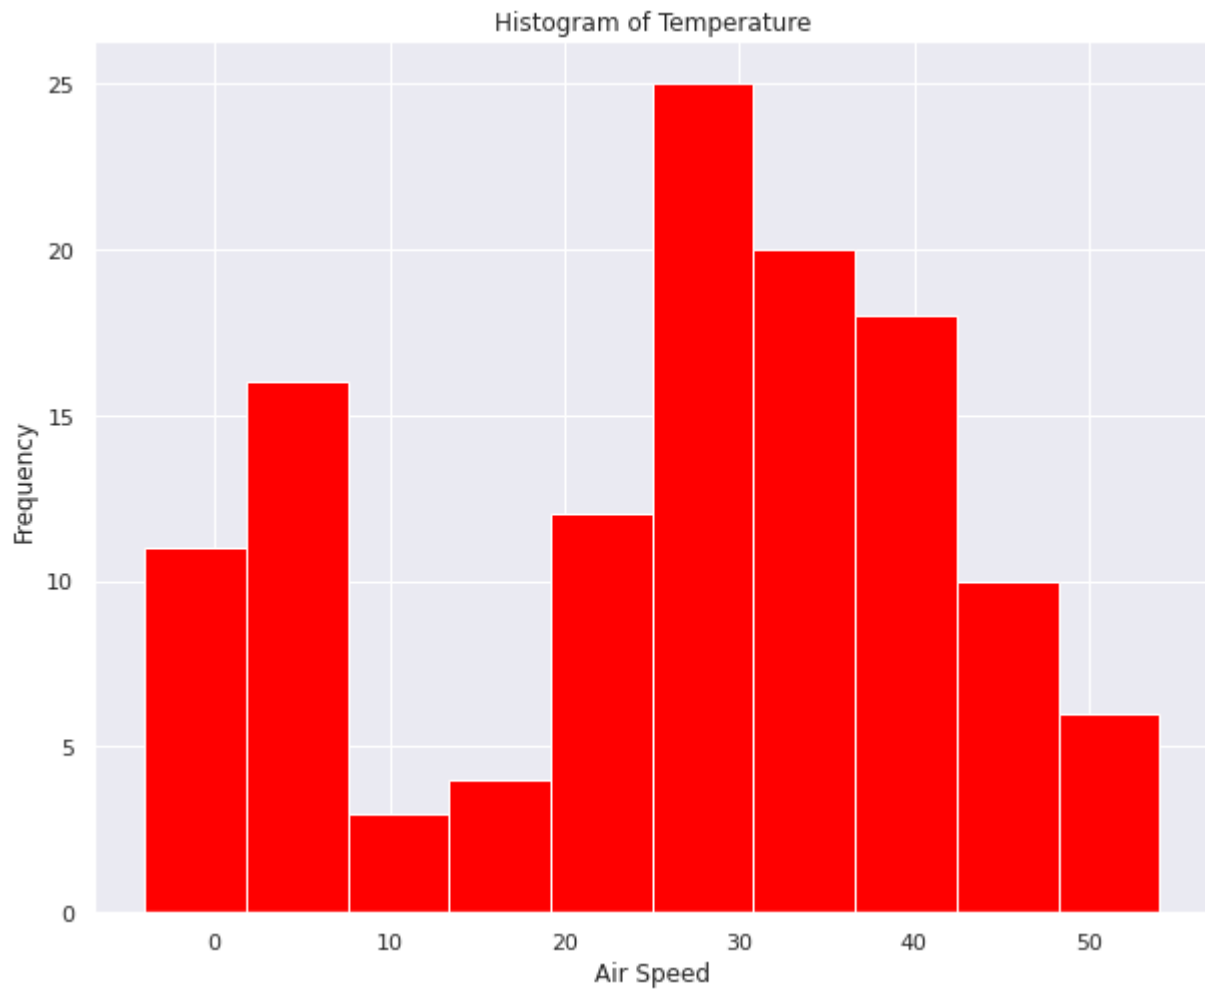
↪ Text(0, 0.5, 'Frequency')



```
# a histogram plot of the temperature column
import matplotlib.pyplot as plt
df_subset.hist('tmpf', color='red')
plt.title("Histogram of Temperature")
plt.xlabel('Air Speed')
plt.ylabel('Frequency')
```

↪

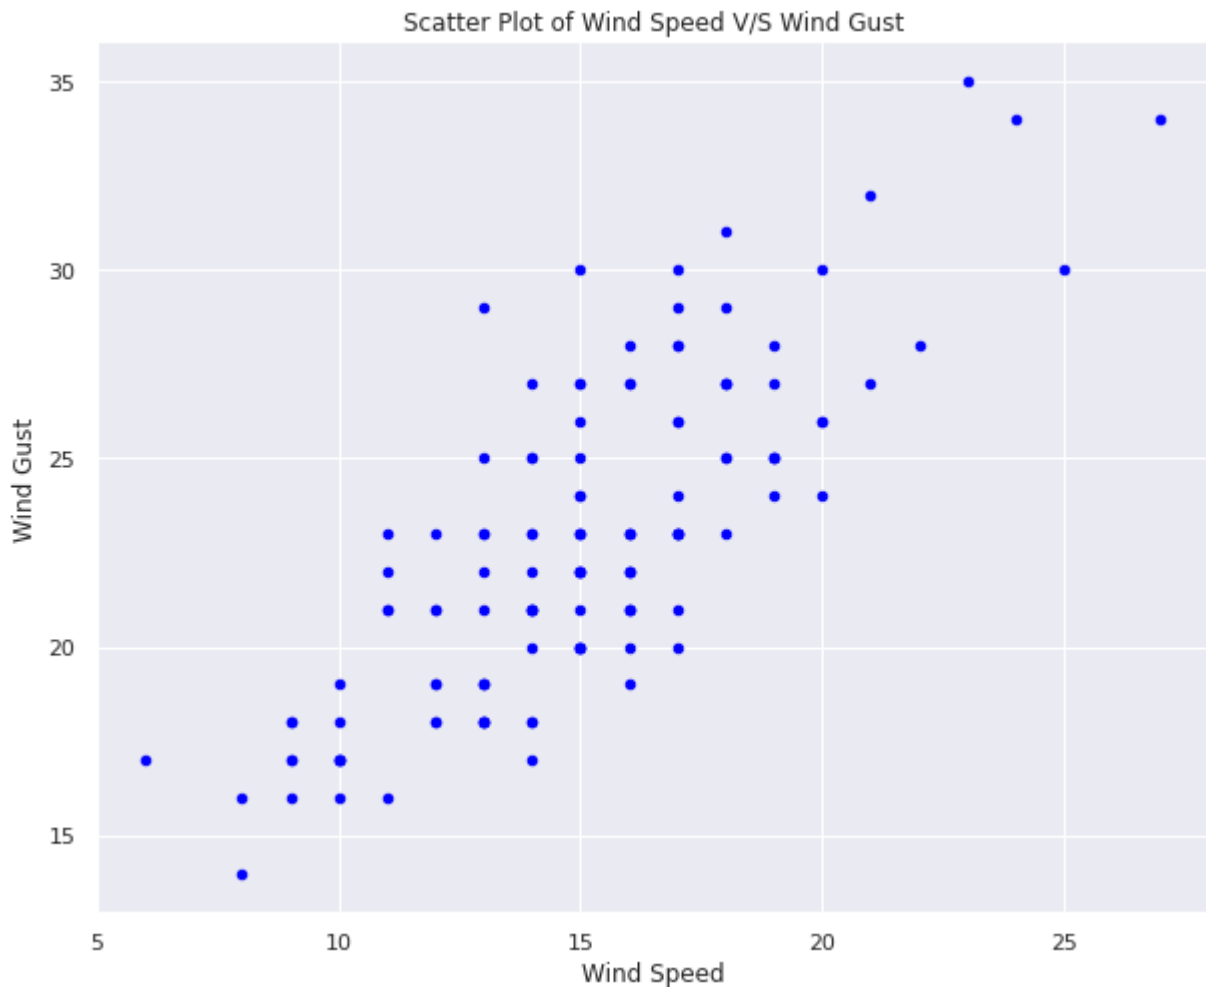
```
Text(0, 0.5, 'Frequency')
```



```
#a scatter plot of the wind speed V/S wind gust column
df_subset['gust'] = df_subset['gust'].astype(float)
df_subset.plot.scatter(x='sknt',
                       y='gust',
                       c='Blue')
plt.title("Scatter Plot of Wind Speed V/S Wind Gust")
plt.xlabel("Wind Speed")
plt.ylabel("Wind Gust")
```




```
Text(0, 0.5, 'Wind Gust')
```



```
#Line plot of the sknt, tmpf and dwpf (as a 3 panel plot), each having a different color.
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
import seaborn as sns
```

```
# specify dimensions of plot
```

```
sns.set(rc={'figure.figsize':(10,8)})
```

```
#3 rows and 1 columns
```

```
fig, ax = plt.subplots(3,1)
```

```
#line plot for Wind Speed
```

```
ax[0].plot(df_subset.sknt,color='green')
```

```
ax[0].set_title("Line Plot of Wind Speed")
```

```
ax[0].set_xlabel('wind speed')
```

```
#line plot for Air Temperature
```

```
ax[1].plot(df_subset.tmpf,color='yellow')
```

```
ax[1].set_title("Line Plot of Air Temperature")
```

```
ax[1].set_xlabel('Air Temperature')
```

```
#line plot for Dew Point Temperature
ax[2].plot(df_subset.dwpf,color='red')
ax[2].set_title("Line Plot of Dew Point Temperature")
ax[2].set_xlabel('Dew Point Temperature')

fig.tight_layout()
plt.show()
```



```
Y_aw = df_subset['gust']
X_aw = df_subset[['tmpf','dwpf','sknt','relh','drct','p01i','alti','mslp']]
X_aw[['dwpf','relh','drct','p01i','alti','mslp']] = X_aw[['dwpf','relh','drct','p01i','alti','mslp']]
X_aw.dtypes
```



```
tmpf      float64
dwpf      float64
sknt      float64
relh      float64
drct      float64
p01i      float64
alti      float64
mslp      float64
dtype: object
```

▼ Polynomial Features

```
# Create polynomial features (and interactions).
from sklearn.preprocessing import PolynomialFeatures

# specify the degree of the poly fit
# the higher the degree, the weirder the fit!
# includes INTERACTIONS and Polynomials... lots of variables here...
poly_features = PolynomialFeatures(degree=2)

# transforms the existing features to higher degree features.
X_poly = poly_features.fit_transform(X_aw)
X_poly = np.nan_to_num(X_poly)
tmp_trans = pd.DataFrame(X_poly)

print("Original Data")
print(X_aw.head()) #13 columns
print("#####")
#tmp_trans= pd.DataFrame(np.nan_to_num(tmp_trans))
print("Poly Features")
print(tmp_trans.head()) #105 columns
tmp_trans.shape
```



Original Data

▼ Log Features

```

# Create log features
import numpy as np
tmp_log_a = np.log(X_aw)
tmp_log_a[np.isneginf(tmp_log_a)] = 0    #recode -inf to 0...
tmp_log = pd.DataFrame(tmp_log_a)
tmp_log = tmp_log.fillna(0)
print(tmp_log.shape) # same number of columns, but they've changed.
tmp_log

```

(125, 8)

	tmpf	dwpf	sknt	relh	drct	p01i	alti	mslp
207	3.972177	3.713572	2.833213	4.149148	5.768321	0.0	3.390473	6.912942
220	3.933784	3.499533	2.944439	3.911623	5.736572	0.0	3.392493	6.915029
232	3.889777	3.430756	2.995732	3.904595	5.736572	0.0	3.393501	6.915823
244	3.828641	3.401197	2.833213	3.977249	5.799093	0.0	3.394844	6.917111
255	3.761200	3.363842	3.295837	4.046729	5.703782	0.0	3.397189	6.919585
...
8141	3.786460	3.295837	2.639057	3.926517	5.247024	0.0	3.397858	6.920375
8297	3.182212	2.302585	2.564949	3.995629	5.828946	0.0	3.403195	6.925595
8950	3.218876	2.714695	2.564949	4.183271	5.799093	0.0	3.394508	6.917210
9053	3.295837	1.945910	2.639057	3.740285	5.438079	0.0	3.393501	6.916021
9077	3.332205	1.945910	2.079442	3.699077	5.247024	0.0	3.392157	6.914532

125 rows × 8 columns

▼ Concatenate log and polynomial features

```

# smoosh the two dataframes together
# reset index of tables
#X_aw.reset_index(inplace=True,drop=True)    #drop not to make additional column for index
tmp_trans.reset_index(inplace=True,drop=True)
tmp_log.reset_index(inplace=True,drop=True)

#Join these two datasets together with the original variables.
df_resultant = pd.concat([tmp_trans,tmp_log], axis=1)

# look at how many more columns there are

```

```
df_resultant.shape
```

```
(125, 53)
```

```
#check missing values
```

```
df_resultant
```

```
missing_values = df_resultant.isnull().sum(axis=0)
```

```
missing_values
```

```
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
19     0
20     0
21     0
22     0
23     0
24     0
25     0
26     0
27     0
28     0
29     0
30     0
31     0
32     0
33     0
34     0
35     0
36     0
37     0
38     0
39     0
40     0
41     0
42     0
43     0
44     0
tmpf   0
dwpf   0
sknt   0
relh   0
drct   0
p01i   0
alti   0
mslp   0
dtvne: int64
```

```
# Shuffle the data, then do an 80/20 split on the data
```

```

Y_aw.reset_index(inplace=True,drop=True)
X_train ,X_test, y_train, y_test = train_test_split(df_resultant, Y_aw, test_size=0.2, random

# ensure the split worked - checked how many rows there are!
print('X_train', X_train.shape,'y_train', y_train.shape)
print('X_test ', X_test.shape, 'y_test ', y_test.shape)

↳ X_train (100, 53) y_train (100,)
   X_test  (25, 53) y_test  (25,)

```

▼ Linear Regression

```

#Linear Regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error

# fit the training data to Linear Regression
LR_model = LinearRegression() # specifies an empty model
LR_model.fit(X_train, y_train)
# predicting on training data-set
y_train_predict = LR_model.predict(X_train)
# predicting on test data-set
y_test_predict = LR_model.predict(X_test)
df_linear_reg_pred = pd.DataFrame(y_test_predict)
df_linear_reg_pred

# the Linear regression model results
score_test = LR_model.score(X_test, y_test)
mae_test = mean_absolute_error(y_test,y_test_predict)
print("R2:{0:.4f}, MAE:{1:.2f}".format(score_test, mae_test))

↳ R2:-0.2335, MAE:3.38

```

▼ Lasso Regression

```

#Fitting an lasso on your concatenated dataset (the polynomial features and
#interactions, and the log features)
#with varying alpha values (0.01, 0.1, 1, 10, 100, 1000),
#observing model best fit in terms of R2 and MAE
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_absolute_error

df_model_results = []

alphas = [0.01, 0.1, 1, 10, 100, 1000]          # this is essentially alpha values

# the Lasso on Training data set

```

```

for a in alphas:
    model = Lasso(alpha=a).fit(X_train, y_train)
    score_train = model.score(X_train, y_train)
    pred_y = model.predict(X_train)
    mae = mean_absolute_error(y_train, pred_y)

# the Lasso on Test data set
score_test = model.score(X_test, y_test)
pred_y_test = model.predict(X_test)
mae_test = mean_absolute_error(y_test, pred_y_test)

#Show results in a table for training and validation results.
df_model_results.append([a, score_train, score_test, mae, mae_test])
df_model_results = pd.DataFrame(data=df_model_results, columns=['Alpha', 'Training R2', 'Test
print(df_model_results)

```

```

↳

```

	Alpha	Training R2	Test R2	Training MAE	Test MAE
0	0.01	0.708028	0.738219	1.876754	1.630231
1	0.10	0.702573	0.747352	1.912473	1.628341
2	1.00	0.698641	0.748369	1.921989	1.615872
3	10.00	0.690043	0.758812	1.973353	1.573544
4	100.00	0.659944	0.756681	2.110637	1.479246
5	1000.00	0.622187	0.703419	2.245159	1.758379

▼ Ridge Regression

```

#Fitting an ridge on your concatenated dataset (the polynomial features and
#interactions, and the log features)
#with varying alpha values (0.01, 0.1, 1, 10, 100, 1000),
#observing model best fit in terms of R2 and MAE
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_absolute_error

df_model_results = []

alphas = [0.01, 0.1, 1, 10, 100, 1000]          # this is essentially alpha values

# the ridge on Training data set
for a in alphas:
    model = Lasso(alpha=a).fit(X_train, y_train)
    score_train = model.score(X_train, y_train)
    pred_y = model.predict(X_train)
    mae = mean_absolute_error(y_train, pred_y)

# the Lasso on Test data set
score_test = model.score(X_test, y_test)
pred_y_test = model.predict(X_test)
mae_test = mean_absolute_error(y_test, pred_y_test)

```

```

#Show results in a table for training and validation results.

```



```
df_model_results.append([a, score_train, score_test, mae, mae_test])
df_model_results = pd.DataFrame(data=df_model_results, columns=['Alpha', 'Training R2', 'Test
print(df_model_results)
```

```
↵
```

	Alpha	Training R2	Test R2	Training MAE	Test MAE
0	0.01	0.708028	0.738219	1.876754	1.630231
1	0.10	0.702573	0.747352	1.912473	1.628341
2	1.00	0.698641	0.748369	1.921989	1.615872
3	10.00	0.690043	0.758812	1.973353	1.573544
4	100.00	0.659944	0.756681	2.110637	1.479246
5	1000.00	0.622187	0.703419	2.245159	1.758379

▼ ElasticNet Regression

```
#Fitting an elasticNet on your concatenated dataset (the polynomial features and
#interactions, and the log features)
#with varying alpha values (0.01, 0.1, 1, 10, 100, 1000), observing model best fit in terms o
from sklearn.linear_model import ElasticNet
from sklearn.metrics import mean_absolute_error
```

```
df_model_results = []
```

```
alphas = [0.01, 0.1, 1, 10, 100, 1000] # this is essentially alpha values
```

```
# the ElasticNet on Training data set
for a in alphas:
    model = ElasticNet(alpha=a).fit(X_train, y_train)
    score_train = model.score(X_train, y_train)
    pred_y = model.predict(X_train)
    mae = mean_absolute_error(y_train, pred_y)
```

```
# the ElasticNet on Test data set
score_test = model.score(X_test, y_test)
pred_y_test = model.predict(X_test)
mae_test = mean_absolute_error(y_test, pred_y_test)
```

```
#Show results in a table for training and validation results.
```

```
df_model_results.append([a, score_train, score_test, mae, mae_test])
df_model_results = pd.DataFrame(data=df_model_results, columns=['Alpha', 'Training R2', 'Test
print(df_model_results)
```


```
↵
```

	Alpha	Training R2	Test R2	Training MAE	Test MAE
0	0.01	0.707933	0.737560	1.877783	1.630985
1	0.10	0.702803	0.747741	1.911686	1.626407
2	1.00	0.699798	0.746914	1.916014	1.623505
3	10.00	0.692084	0.754303	1.957291	1.586926
4	100.00	0.676581	0.775164	2.033666	1.464900
5	1000.00	0.626251	0.693646	2.235905	1.776960

```
#Scatterplot of actual vs. predicted for your best fitting model.
```

```
import seaborn
model = ElasticNet(alpha=100.00).fit(X_train, y_train)
score_train = model.score(X_train, y_train)
pred_y = model.predict(X_train)
seaborn.regplot(y_train, pred_y)
plt.title("Actual V/S Predicted Training Data")
plt.xlabel("Actual")
plt.ylabel("Predicted")
```

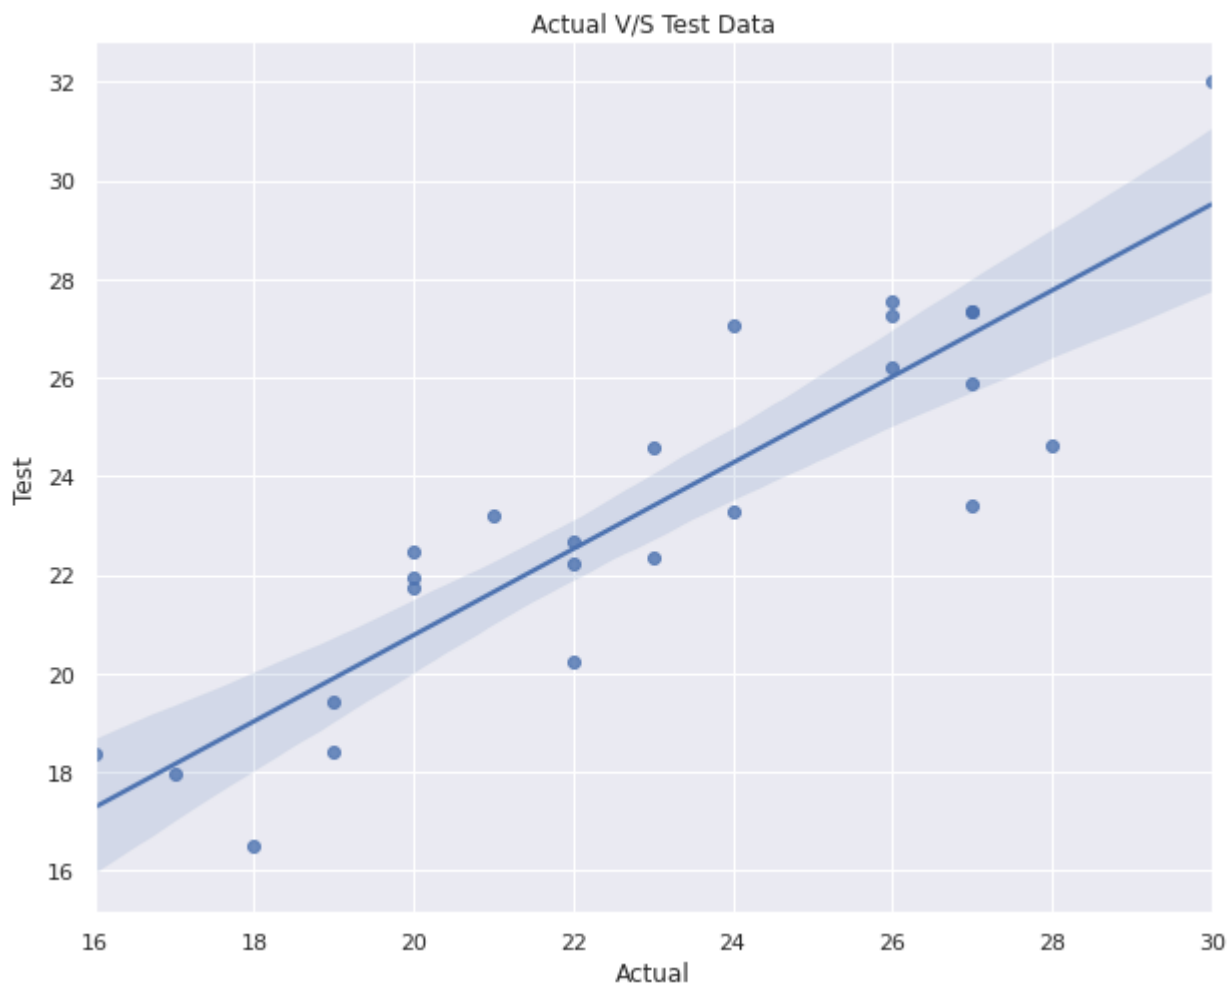
↗ Text(0, 0.5, 'Predicted')



```
#Scatterplot of actual vs. test for your best fitting model.
score_test = model.score(X_test, y_test)
pred_y_test = model.predict(X_test)
mae_test = mean_absolute_error(y_test, pred_y_test)
seaborn.regplot(y_test, pred_y_test)
plt.title("Actual V/S Test Data")
plt.xlabel("Actual")
plt.ylabel("Test")
```

↗

```
Text(0, 0.5, 'Test')
```



▼ Baseline ElasticNet Regression

```
#After doing the same on the smaller original dataset it was observed that
#all the feature engineering helped enhance model accuracy
X_aw.reset_index(inplace=True,drop=True) #drop not to make additional column for index
X_train_subset ,X_test_subset, Y_train_subset, Y_test_subset = train_test_split(X_aw, Y_aw, t

# ensure the split worked - checked how many rows there are!
print('X_train', X_train_subset.shape,'y_train', Y_train_subset.shape)
print('X_test ', X_test_subset.shape, 'y_test ', X_test_subset.shape)
```

```
↳ X_train (100, 8) y_train (100,)
   X_test (25, 8) y_test (25, 8)
```

```
alphas = [0.01, 0.1, 1, 10, 100, 1000]
```

```
df_model_results_s = []
```

```

# the ElasticNet on original dataset
# training dataset
for a in alphas:
    model = ElasticNet(alpha=a).fit(X_train_subset, Y_train_subset)
    score_train = model.score(X_train_subset, Y_train_subset)
    pred_y_simple_train = model.predict(X_train_subset)
    mae = mean_absolute_error(Y_train_subset, pred_y_simple_train)
    score_test = model.score(X_test_subset, Y_test_subset)
    pred_y_simple_test = model.predict(X_test_subset)
    mae_test = mean_absolute_error(Y_test_subset, pred_y_simple_test)

#Show results in a table for training and validation results.
df_model_results_s.append([a, score_train, score_test, mae, mae_test])
df_model_results_s = pd.DataFrame(data=df_model_results_s, columns=['Alpha', 'Training R2', '
print(df_model_results_s)

```

	Alpha	Training R2	Test R2	Training MAE	Test MAE
0	0.01	0.640694	0.672966	2.201683	1.804642
1	0.10	0.638061	0.680525	2.202587	1.793387
2	1.00	0.623941	0.702218	2.240013	1.760866
3	10.00	0.413763	0.525454	2.701517	2.230617
4	100.00	0.000000	-0.003528	3.480000	3.170400
5	1000.00	0.000000	-0.003528	3.480000	3.170400